

Chapter 12. Epilogue: Apache Spark 3.0

At the time we were writing this book, Apache Spark 3.0 had not yet been officially released; it was still under development, and we got to work with Spark 3.0.0-preview2. All the code samples in this book have been tested against Spark 3.0.0-preview2, and they should work no differently with the official Spark 3.0 release. Whenever possible in the chapters, where relevant, we mentioned when features were new additions or behaviors in Spark 3.0. In this chapter, we survey the changes.

The bug fixes and feature enhancements are numerous, so for brevity, we highlight just a selection of the notable changes and features pertaining to Spark components. Some of the new features are, under the hood, advanced and beyond the scope of this book, but we mention them here so you can explore them when the release is generally available.

Spark Core and Spark SQL

Let's first consider what's new under the covers. A number of changes have been introduced in Spark Core and the Spark SQL engine to help speed up queries. One way to expedite queries is to read less data using dynamic partition pruning. Another is to adapt and optimize query plans during execution.

Dynamic Partition Pruning

The idea behind [dynamic partition pruning \(DPP\)](#) is to skip over the data you don't need in a query's results. The typical scenario where DPP is optimal is when you are joining two tables: a fact table (partitioned over multiple columns) and a dimension table (nonpartitioned), as shown in [Figure 12-1](#). Normally, the filter is on the nonpartitioned side of the table (Date, in our case). For example, consider this common query over two tables, Sales and Date:

```
-- In SQL
```

```
SELECT * FROM Sales JOIN ON Sales.date = Date.date
```

Figure 12-1. Dynamic filter is injected from the dimension table into the fact table

The key optimization technique in DPP is to take the result of the filter from the dimension table and inject it into the fact table as part of the scan operation to limit the data read, as shown in [Figure 12-1](#).

Consider a case where the dimension table is smaller than the fact table and we perform a join, as shown in [Figure 12-2](#). In this case, Spark most likely will do a broadcast join (discussed in [Chapter 7](#)). During this join, Spark will conduct the following steps to minimize the amount of data scanned from the larger fact table:

1. On the dimension side of the join, Spark will build a hash table from the dimension table, also known as the build relation, as part of this filter query.
2. Spark will plug the result of this query into the hash table and assign it to a broadcast variable, which is distributed to all executors involved in this join operation.
3. On each executor, Spark will probe the broadcasted hash table to determine what corresponding rows to read from the fact table.
4. Finally, Spark will inject this filter dynamically into the file scan operation of the fact table and reuse the results from the broadcast variable. This way, as part of the file scan operation on the fact table, only the partitions that match the filter are scanned and only the data needed is read.

Figure 12-2. Spark injects a dimension table filter into the fact table during a broadcast join

Enabled by default so that you don't have to explicitly configure it, all this happens dynamically when you perform joins between two tables. With the DPP optimization, Spark 3.0 can work much better with star-schema queries.

Adaptive Query Execution

Another way Spark 3.0 optimizes query performance is by adapting its physical execution plan at runtime. [Adaptive Query Execution \(AQE\)](#) reoptimizes and adjusts query plans based on runtime statistics collected in the process of query execution. It attempts to do the following at runtime:

- Reduce the number of reducers in the shuffle stage by decreasing the number of shuffle partitions.
- Optimize the physical execution plan of the query, for example by converting a `SortMergeJoin` into a `BroadcastHashJoin` where appropriate.
- Handle data skew during a join.

All these adaptive measures take place during the execution of the plan at runtime, as shown in [Figure 12-3](#). To use AQE in Spark 3.0, set the configuration `spark.sql.adaptive.enabled` to `true`.

Figure 12-3. AQE reexamines and reoptimizes the execution plan at runtime

The AQE framework

Spark operations in a query are pipelined and executed in parallel processes, but a shuffle or broadcast exchange breaks this pipeline, because the output of one stage is needed as input to the next stage (see [“Step 3: Understanding Spark Application Concepts”](#) in [Chapter 2](#)). These breaking points are called *materialization points* in a query stage, and they present an opportunity to reoptimize and reexamine the query, as illustrated in [Figure 12-4](#).

Figure 12-4. A query plan reoptimized in the AQE framework

Here are the conceptual steps the AQE framework iterates over, as depicted in this figure:

1. All the leaf nodes, such as scan operations, of each stage are executed.

2. Once the materialization point finishes executing, it's marked as complete, and all the relevant statistics garnered during execution are updated in its logical plan.
3. Based on these statistics, such as number of partitions read, bytes of data read, etc., the framework runs the Catalyst optimizer again to understand whether it can:
 1. Coalesce the number of partitions to reduce the number of reducers to read shuffle data.
 2. Replace a sort merge join, based on the size of tables read, with a broadcast join.
 3. Try to remedy a skew join.
 4. Create a new optimized logical plan, followed by a new optimized physical plan.

This process is repeated until all the stages of the query plan are executed.

In short, this reoptimization is done dynamically, as shown in [Figure 12-3](#), and the objective is to dynamically coalesce the shuffle partitions, decrease the number of reducers needed to read the shuffle output data, switch join strategies if appropriate, and remedy any skew joins.

Two Spark SQL configurations dictate how AQE will reduce the number of reducers:

- `spark.sql.adaptive.coalescePartitions.enabled` (set to `true`)
- `spark.sql.adaptive.skewJoin.enabled` (set to `true`)

At the time of writing, the Spark 3.0 community blog, documentation, and examples had not been published publicly, but by the time of publication they should have been. These resources will enable you to get more detailed information if you wish to see how these features work under the hood—including on how you can inject SQL join hints, discussed next.

SQL Join Hints

Adding to the existing `BROADCAST` hints for joins, Spark 3.0 adds join hints for all [Spark join strategies](#) (see [“A Family of Spark Joins”](#) in [Chapter 7](#)). Examples are provided here for each type of join.

Shuffle sort merge join (SMJ)

With these new hints, you can suggest to Spark that it perform a `SortMergeJoin` when joining tables `a` and `b` or `customers` and `orders`, as shown in the following examples. You can add one or more hints to a `SELECT` statement inside `/*+ ... */` comment blocks:

```
SELECT /*+ MERGE(a, b) */ id FROM a JOIN b ON a.key = b.key
SELECT /*+ MERGE(customers, orders) */ * FROM customers, orders WHERE
    orders.custId = customers.custId
```

Broadcast hash join (BHJ)

Similarly, for a broadcast hash join, you can provide a hint to Spark that you prefer a broadcast join. For example, here we broadcast table `a` to join with table `b` and table `customers` to join with table `orders`:

```
SELECT /*+ BROADCAST(a) */ id FROM a JOIN b ON a.key = b.key
SELECT /*+ BROADCAST(customers) */ * FROM customers, orders WHERE
    orders.custId = customers.custId
```

Shuffle hash join (SHJ)

You can offer hints in a similar way to perform shuffle hash joins, though this is less commonly encountered than the previous two supported join strategies:

```
SELECT /*+ SHUFFLE_HASH(a, b) */ id FROM a JOIN b ON a.key = b.key
SELECT /*+ SHUFFLE_HASH(customers, orders) */ * FROM customers, orders WHERE
    orders.custId = customers.custId
```

Shuffle-and-replicate nested loop join (SNLJ)

Finally, the shuffle-and-replicate nested loop join adheres to a similar form and syntax:

```
SELECT /*+ SHUFFLE_REPLICATE_NL(a, b) */ id FROM a JOIN b
```

Catalog Plugin API and DataSourceV2

Not to be confined only to the Hive metastore and catalog, Spark 3.0's experimental DataSourceV2 API extends the Spark ecosystem and affords developers three core capabilities. Specifically, it:

- Enables plugging in an external data source for catalog and table management
- Supports predicate pushdown to additional data sources with supported file formats like ORC, Parquet, Kafka, Cassandra, Delta Lake, and Apache Iceberg.
- Provides unified APIs for streaming and batch processing of data sources for sinks and sources

Aimed at developers who want to extend Spark's ability to use external sources and sinks, the Catalog API provides both SQL and programmatic APIs to create, alter, load, and drop tables from the specified pluggable catalog. The catalog provides a hierarchical abstraction of functionalities and operations performed at different levels, as shown in [Figure 12-5](#).

Figure 12-5. Catalog plugin API's hierarchical level of functionality

The initial interaction between Spark and a specific connector is to resolve a relation to its actual Table object. Catalog defines how to look up tables in this connector. Additionally, Catalog can define how to modify its own metadata, thus [enabling operations](#) like CREATE TABLE , ALTER TABLE , etc.

For example, in SQL you can now issue commands to create namespaces for your catalog. To use a pluggable catalog, enable the following configs in your *spark-defaults.conf* file:

```
spark.sql.catalog.ndb_catalog com.ndb.ConnectorImpl # connector implementation
spark.sql.catalog.ndb_catalog.option1 value1
spark.sql.catalog.ndb_catalog.option2 value2
```

Here, the connector to the data source catalog has two options: option1->value1 and option2->value2 . Once they've been defined, application

users in Spark or SQL can use the `DataFrameReader` and `DataFrameWriter` API methods or Spark SQL commands with these defined options as methods for data source manipulation. For example:

```
-- In SQL
SHOW TABLES ndb_catalog;
CREATE TABLE ndb_catalog.table_1;
SELECT * from ndb_catalog.table_1;
ALTER TABLE ndb_catalog.table_1

// In Scala
df.writeTo("ndb_catalog.table_1")
val dfNBD = spark.read.table("ndb_catalog.table_1")
    .option("option1", "value1")
    .option("option2", "value2")
```

While these catalog plugin APIs extend Spark's ability to utilize external data sources as sinks and sources, they are still experimental and should not be used in production. A detailed guide to their use is beyond the scope of this book, but we encourage you to check the release documentation for additional information if you want to write a custom connector to an external data source as a catalog to manage your external tables and their associated metadata.

NOTE

The preceding code snippets are examples of what your code may look like after you have defined and implemented your catalog connectors and populated them with data.

Accelerator-Aware Scheduler

[Project Hydrogen](#), a community initiative to bring AI and big data together, has three major goals: implementing barrier execution mode, accelerator-aware scheduling, and optimized data exchange. A basic implementation of [barrier execution mode](#) was introduced in Apache Spark 2.4.0. In Spark 3.0, a basic [scheduler](#) has been implemented to take advan-

tage of hardware accelerators such as GPUs on target platforms where Spark is deployed in standalone mode, YARN, or Kubernetes.

For Spark to take advantage of these GPUs in an organized way for specialized workloads that use them, you have to specify the hardware resources available via configs. Your application can then discover them with the help of a discovery script. Enabling GPU use is a three-step process in your Spark application:

1. Write a discovery script that discovers the addresses of the underlying GPUs available on each Spark executor. This script is set in the following Spark configuration:

```
spark.worker.resource.gpu.discoveryScript=/path/to/script.sh
```

2. Set up configuration for your Spark executors to use these discovered GPUs:

```
spark.executor.resource.gpu.amount=2
spark.task.resource.gpu.amount=1
```

3. Write RDD code to leverage these GPUs for your task:

```
import org.apache.spark.BarrierTaskContext
val rdd = ...
rdd.barrier.mapPartitions { it =>
    val context = BarrierTaskContext.getcontext.barrier()
    val gpus = context.resources().get("gpu").get.addresses
    // launch external process that leverages GPU
    launchProcess(gpus)
}
```

NOTE

These steps are still experimental, and further development will continue in future Spark 3.x releases to support seamless discovery of GPU resources, both at the command line (with `spark-submit`) and at the Spark task level.

Structured Streaming

To inspect how your Structured Streaming jobs fare with the ebb and flow of data during the course of execution, the Spark 3.0 UI has a new Structured Streaming tab alongside the other tabs we explored in [Chapter 7](#). This tab offers two sets of statistics: aggregate information about completed streaming query jobs ([Figure 12-6](#)) and detailed statistics about the streaming queries, including the input rate, process rate, number of input rows, batch duration, and operation duration ([Figure 12-7](#)).

Figure 12-6. Structured Streaming tab showing aggregate statistics of a completed streaming job

NOTE

The [Figure 12-7](#) screenshot was taken with Spark 3.0.0-preview2; with the final release, you should see the query name and ID in the name identifier on the UI page.

Figure 12-7. Showing detailed statistics of a completed streaming job

[No configuration is required](#); all configurations works straight out of the Spark 3.0 installation, with the following defaults:

- `spark.sql.streaming.ui.enabled=true`
- `spark.sql.streaming.ui.retainedProgressUpdates=100`
- `spark.sql.streaming.ui.retainedQueries=100`

PySpark, Pandas UDFs, and Pandas Function APIs

Spark 3.0 requires pandas version 0.23.2 or higher to employ any pandas-related methods, such as `DataFrame.toPandas()` or `SparkSession.createDataFrame(pandas.DataFrame)`.

Furthermore, it requires PyArrow version 0.12.1 or later to use PyArrow functionality such as `pandas_udf()`, `DataFrame.toPandas()`, and `SparkSession.createDataFrame(pandas.DataFrame)` with the `spark.sql.execution.arrow.enabled` configuration set to `true`. The next section will introduce new features in Pandas UDFs.

Redesigned Pandas UDFs with Python Type Hints

The Pandas UDFs in Spark 3.0 were redesigned by leveraging [Python type hints](#). This enables you to naturally express UDFs without requiring the evaluation type. Pandas UDFs are now more “Pythonic” and can themselves define what the UDF is supposed to input and output, rather than you specifying it via, for example, `@pandas_udf("long", PandasUDFType.SCALAR)` as you did in Spark 2.4.

Here’s an example:

```
# Pandas UDFs in Spark 3.0
import pandas as pd
from pyspark.sql.functions import pandas_udf

@pandas_udf("long")
def pandas_plus_one(v: pd.Series) -> pd.Series:
    return v + 1
```

This new format provides several benefits, such as easier static analysis. You can apply the new UDFs in the same way as before:

```
df = spark.range(3)
df.withColumn("plus_one", pandas_plus_one("id")).show()
```

```
+---+-----+
| id|plus_one|
+---+-----+
|  0|        1|
|  1|        2|
|  2|        3|
+---+-----+
```

Iterator Support in Pandas UDFs

Pandas UDFs are very commonly used to load a model and perform distributed inference for single-node machine learning and deep learning models. However, if a model is very large, then there is high overhead for the Pandas UDF to repeatedly load the same model for every batch in the same Python worker process.

In Spark 3.0, Pandas UDFs can [accept an iterator](#) of `pandas.Series` or `pandas.DataFrame`, as shown here:

```
from typing import Iterator

@pandas_udf('long')
def pandas_plus_one(iterator: Iterator[pd.Series]) -> Iterator[pd.Series]:
    return map(lambda s: s + 1, iterator)

df.withColumn("plus_one", pandas_plus_one("id")).show()

+---+-----+
| id|plus_one|
+---+-----+
|  0|         1|
|  1|         2|
|  2|         3|
+---+-----+
```

With this support, you can load the model only once instead of loading it for every series in the iterator. The following pseudocode illustrates how to do this:

```
@pandas_udf(...)
def predict(iterator):
    model = ... # load model
    for features in iterator:
        yield model.predict(features)
```

New Pandas Function APIs

Spark 3.0 introduces a few new types of Pandas UDFs that are useful when you want to apply a function against an entire DataFrame instead of column-wise, such as `mapInPandas()`, introduced in [Chapter 11](#). These take an iterator of `pandas.DataFrame` as input and output another iterator of `pandas.DataFrame`:

```
def pandas_filter(
    iterator: Iterator[pd.DataFrame]) -> Iterator[pd.DataFrame]:
    for pdf in iterator:
        yield pdf[pdf.id == 1]
```

```
df.mapInPandas(pandas_filter, schema=df.schema).show()
```

```
+----+
| id |
+----+
|  1 |
+----+
```

You can control the size of the `pandas.DataFrame` by specifying it in the `spark.sql.execution.arrow.maxRecordsPerBatch` configuration. Note that the input size and output size do not have to match, unlike with most Pandas UDFs.

NOTE

All the data of a cogroup will be loaded into memory, which means if there is data skew or certain groups are too big to fit in memory you could run into OOM issues.

Spark 3.0 also introduces cogrouped map Pandas UDFs. The `applyInPandas()` function takes two `pandas.DataFrame`s that share a common key and applies a function to each cogroup. The returned `pandas.DataFrame`s are then combined as a single DataFrame. As with `mapInPandas()`, there is no restriction on the length of the returned `pandas.DataFrame`. Here's an example:

```

df1 = spark.createDataFrame(
    [(1201, 1, 1.0), (1201, 2, 2.0), (1202, 1, 3.0), (1202, 2, 4.0)],
    ("time", "id", "v1"))
df2 = spark.createDataFrame(
    [(1201, 1, "x"), (1201, 2, "y")], ("time", "id", "v2"))

def asof_join(left: pd.DataFrame, right: pd.DataFrame) -> pd.DataFrame:
    return pd.merge_asof(left, right, on="time", by="id")

df1.groupby("id").cogroup(
    df2.groupby("id")
).applyInPandas(asof_join, "time int, id int, v1 double, v2 string").show()

+----+----+----+----+
|time| id| v1| v2|
+----+----+----+----+
|1201|  1|1.0|  x|
|1202|  1|3.0|  x|
|1201|  2|2.0|  y|
|1202|  2|4.0|  y|
+----+----+----+----+

```

Changed Functionality

Listing all the functionality changes in Spark 3.0 would transform this book into a brick several inches thick. So, in the interest of brevity, we will mention a few notable ones here, and leave you to consult the release notes for Spark 3.0 for full details and all the nuances as soon as they are available.

Languages Supported and Deprecated

Spark 3.0 supports Python 3 and JDK 11, and Scala version 2.12 is required. All Python versions earlier than 3.6 and Java 8 are deprecated. If you use these deprecated versions you will get warning messages.

Changes to the DataFrame and Dataset APIs

In previous versions of Spark, the Dataset and DataFrame APIs had deprecated the `unionAll()` method. In Spark 3.0 this has been reversed, and

`unionAll()` is now an alias to the `union()` method.

Also, earlier versions of Spark's `Dataset.groupByKey()` resulted in a grouped Dataset with the key spuriously named as `value` when the key was a non-struct type (`int`, `string`, `array`, etc.). As such, aggregation results from `ds.groupByKey().count()` in the query when displayed looked, counterintuitively, like `(value, count)`. This has been rectified to result in `(key, count)`, which is more intuitive. For example:

```
// In Scala
val ds = spark.createDataset(Seq(20, 3, 3, 2, 4, 8, 1, 1, 3))
ds.show(5)
```

```
+-----+
|value|
+-----+
|  20|
|   3|
|   3|
|   2|
|   4|
+-----+
```

```
ds.groupByKey(k=> k).count.show(5)
```

```
+---+-----+
|key|count(1)|
+---+-----+
|  1|        2|
|  3|        3|
| 20|        1|
|  4|        1|
|  8|        1|
+---+-----+
```

```
only showing top 5 rows
```

However, you can preserve the old format if you prefer by setting `spark.sql.legacy.dataset.nameNonStructGroupingKeyAsValue` to `true`.

DataFrame and SQL Explain Commands

For better readability and formatting, Spark 3.0 introduces the `DataFrame.explain(FORMAT_MODE)` capability to display different views of the plans the Catalyst optimizer generates. The `FORMAT_MODE` options include "simple" (the default), "extended", "cost", "codegen", and "formatted". Here's a simple illustration:

```
// In Scala
val strings = spark
  .read.text("/databricks-datasets/learning-spark-v2/SPARK_README.md")
val filtered = strings.filter($"value".contains("Spark"))
filtered.count()
```

```
# In Python
strings = spark
  .read.text("/databricks-datasets/learning-spark-v2/SPARK_README.md")
filtered = strings.filter(strings.value.contains("Spark"))
filtered.count()
```

```
// In Scala
filtered.explain("simple")
```

```
# In Python
filtered.explain(mode="simple")
```

```
== Physical Plan ==
*(1) Project [value#72]
+- *(1) Filter (isnotnull(value#72) AND Contains(value#72, Spark))
   +- FileScan text [value#72] Batched: false, DataFilters: [isnotnull(value#72),
Contains(value#72, Spark)], Format: Text, Location:
InMemoryFileIndex[dbfs:/databricks-datasets/learning-spark-v2/SPARK_README.md],
PartitionFilters: [], PushedFilters: [IsNotNull(value),
StringContains(value,Spark)], ReadSchema: struct<value:string>
```

```
// In Scala
filtered.explain("formatted")
```

```

# In Python
filtered.explain(mode="formatted")

== Physical Plan ==
* Project (3)
+- * Filter (2)
   +- Scan text (1)

(1) Scan text
Output [1]: [value#72]
Batched: false
Location: InMemoryFileIndex [dbfs:/databricks-datasets/learning-spark-v2/...
PushedFilters: [IsNotNull(value), StringContains(value,Spark)]
ReadSchema: struct<value:string>

(2) Filter [codegen id : 1]
Input [1]: [value#72]
Condition : (isnotnull(value#72) AND Contains(value#72, Spark))

(3) Project [codegen id : 1]
Output [1]: [value#72]
Input [1]: [value#72]

-- In SQL
EXPLAIN FORMATTED
SELECT *
FROM tmp_spark_readme
WHERE value like "%Spark%"

== Physical Plan ==
* Project (3)
+- * Filter (2)
   +- Scan text (1)

(1) Scan text
Output [1]: [value#2016]
Batched: false
Location: InMemoryFileIndex [dbfs:/databricks-datasets/
learning-spark-v2/SPARK_README.md]
PushedFilters: [IsNotNull(value), StringContains(value,Spark)]
ReadSchema: struct<value:string>

(2) Filter [codegen id : 1]

```



```
Input [1]: [value#2016]
Condition : (isnotnull(value#2016) AND Contains(value#2016, Spark))

(3) Project [codegen id : 1]
Output [1]: [value#2016]
Input [1]: [value#2016]
```

To see the rest of the format modes in action, you can try the notebook in the book's [GitHub repo](#). Also check out the [migration guides](#) from Spark 2.x to Spark 3.0.

Summary

This chapter provided a cursory highlight of new features in Spark 3.0. We took the liberty of mentioning a few advanced features that are worthy of note. They operate under the hood and not at the API level. In particular, we took a look at dynamic partition pruning (DPP) and adaptive query execution (AQE), two optimizations that enhance Spark's performance at execution time. We also explored how the experimental Catalog API extends the Spark ecosystem to custom data stores for sources and sinks for both batch and streaming data, and looked at the new scheduler in Spark 3.0 that enables it to take advantage of GPUs in executors.

Complementing our discussion of the Spark UI in [Chapter 7](#), we also showed you the new Structured Streaming tab, providing accumulated statistics on streaming jobs, additional visualizations, and detailed metrics on each query.

Python versions below 3.6 are deprecated in Spark 3.0, and Pandas UDFs have been redesigned to support Python type hints and iterators as arguments. There are Pandas UDFs that enable transforming an entire DataFrame, as well as combining two cogrouped DataFrames into a new DataFrame.

For better readability of query plans, `DataFrame.explain(FORMAT_MODE)` and `EXPLAIN FORMAT_MODE` in SQL display different levels and details of logical and physical plans. Additionally, SQL commands can now take join hints for Spark's entire supported family of joins.

While we were unable to enumerate all the changes in the latest version of Spark in this short chapter, we urge that you explore the release notes when Spark 3.0 is released to find out more. Also, for a quick summary of the user-facing changes and details on how to migrate to Spark 3.0, we encourage you to check out the migration guides.

As a reminder, all the code in this book has been tested on Spark 3.0.0-preview2 and should work with Spark 3.0 when it is officially released. We hope you've enjoyed reading this book and learned from this journey with us. We thank you for your attention!