


Hadoop Architecture and HDFS



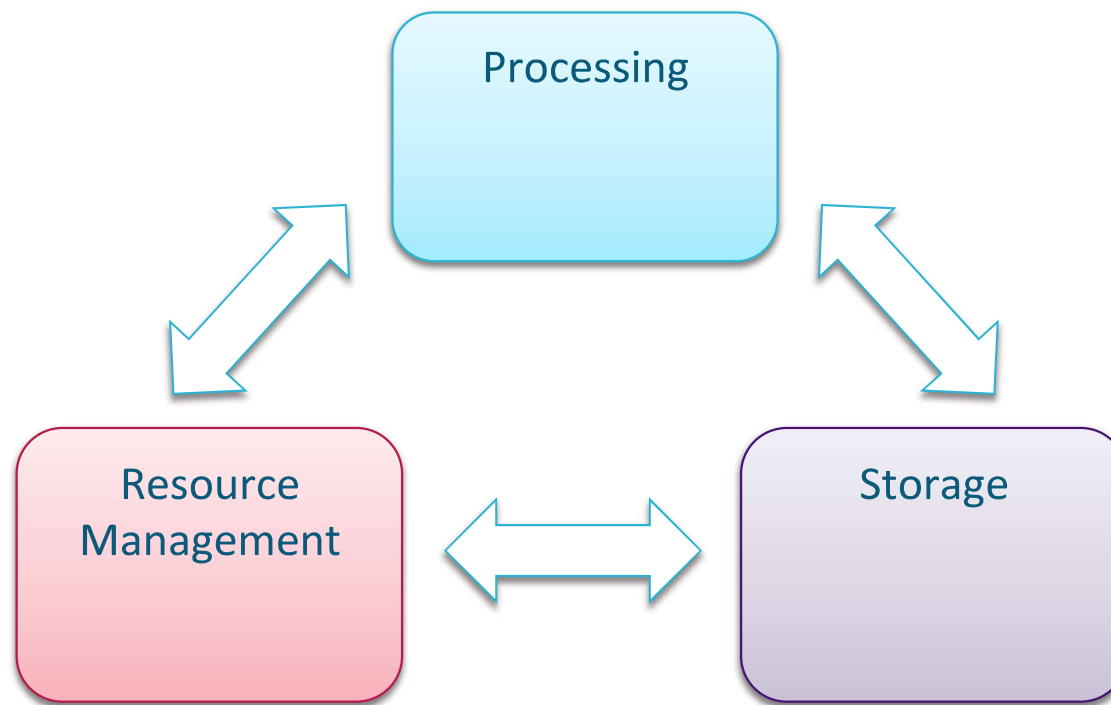
Hadoop Cluster Terminology

- A **cluster** is a group of computers working together
 - Provides data storage, data processing, and resource management
- A **node** is an individual computer in the cluster
 - *Master* nodes manage distribution of work and data to *worker* nodes
- A **daemon** is a program running on a node
 - Each Hadoop daemon performs a specific function in the cluster




Cluster Components

- **Three main components of a cluster**
- **Work together to provide distributed data processing**
- **We will start with the Storage component**
 - HDFS



HDFS Basic Concepts (1)

- **HDFS is a filesystem written in Java**
 - Based on Google's GFS
- **Sits on top of a native filesystem**
 - Such as ext3, ext4, or xfs
- **Provides redundant storage for massive amounts of data**
 - Using readily-available, industry-standard computers

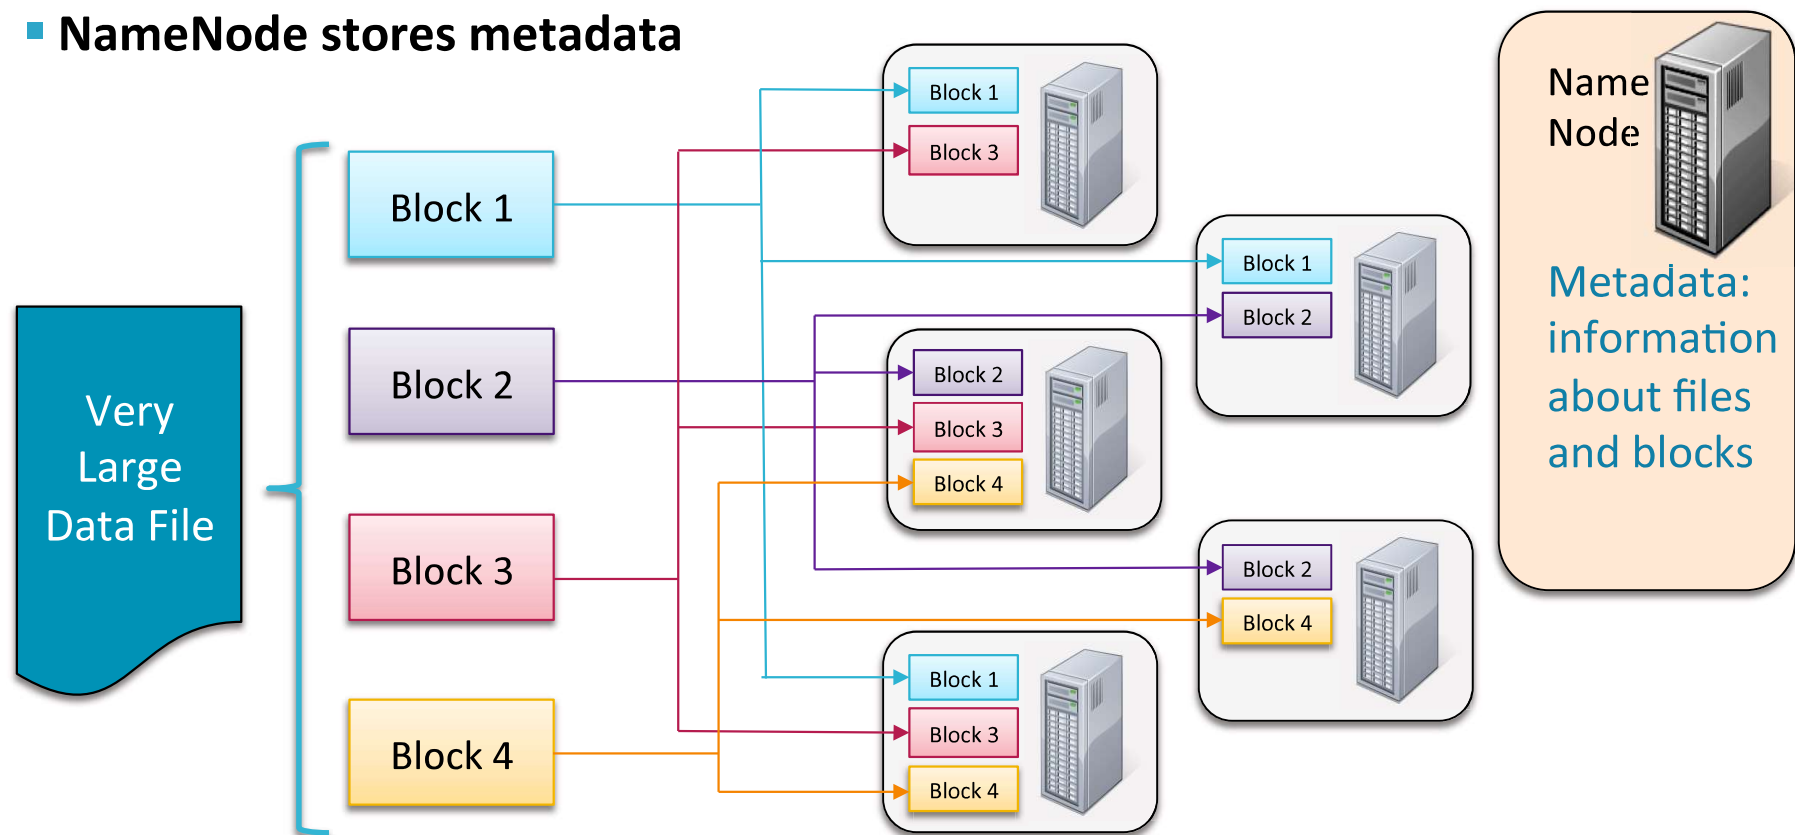


HDFS Basic Concepts (2)

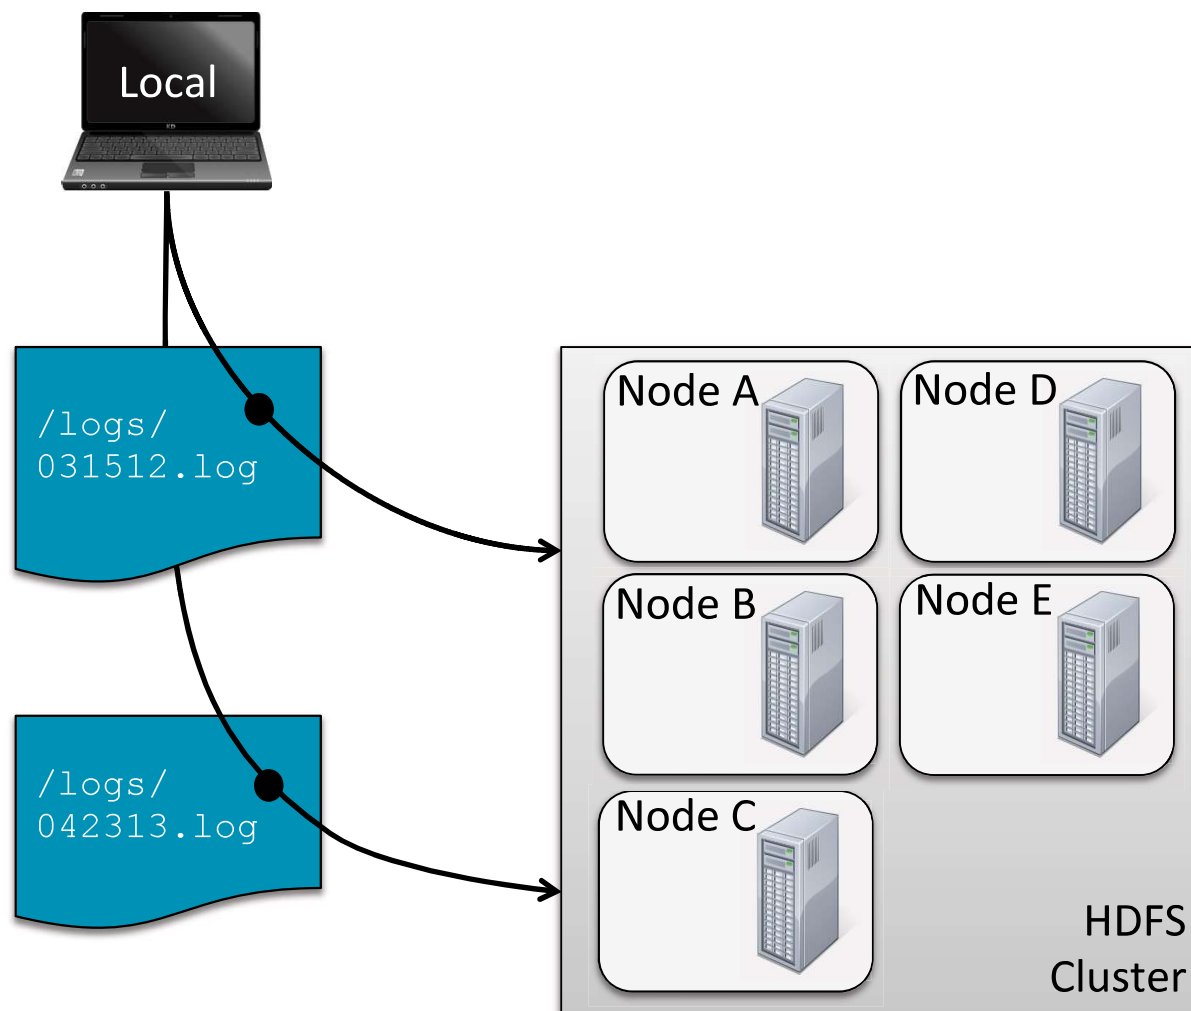
- **HDFS performs best with a ‘modest’ number of large files**
 - Millions, rather than billions, of files
 - Each file typically 100MB or more
- **Files in HDFS are ‘write once’**
 - No random writes to files are allowed
- **HDFS is optimized for large, streaming reads of files**
 - Rather than random reads

How Files Are Stored

- Data files are split into 128MB blocks which are distributed at load time
- Each block is replicated on multiple data nodes (default 3x)
- NameNode stores metadata



Example: Storing and Retrieving Files (1)



Example: Storing and Retrieving Files (2)

Metadata

`/logs/031512.log`: B1, B2, B3
`/logs/042313.log`: B4, B5

B1: A, B, D

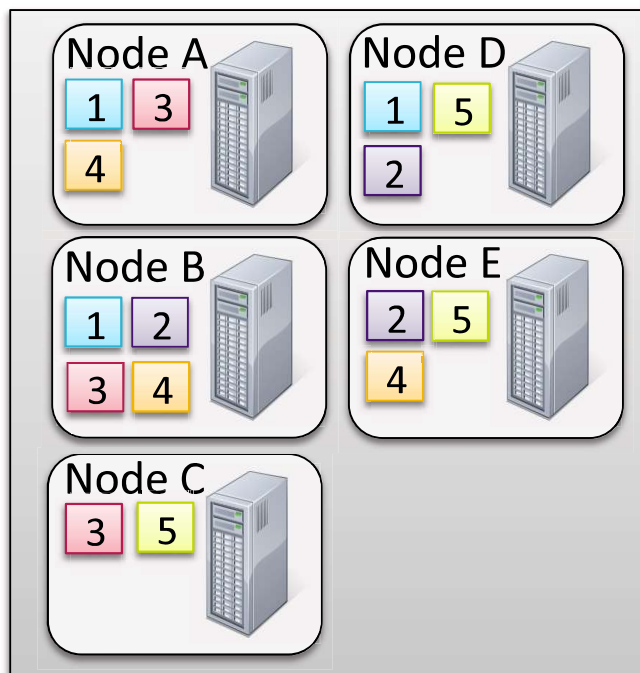
B2: B, D, E

B3: A, B, C

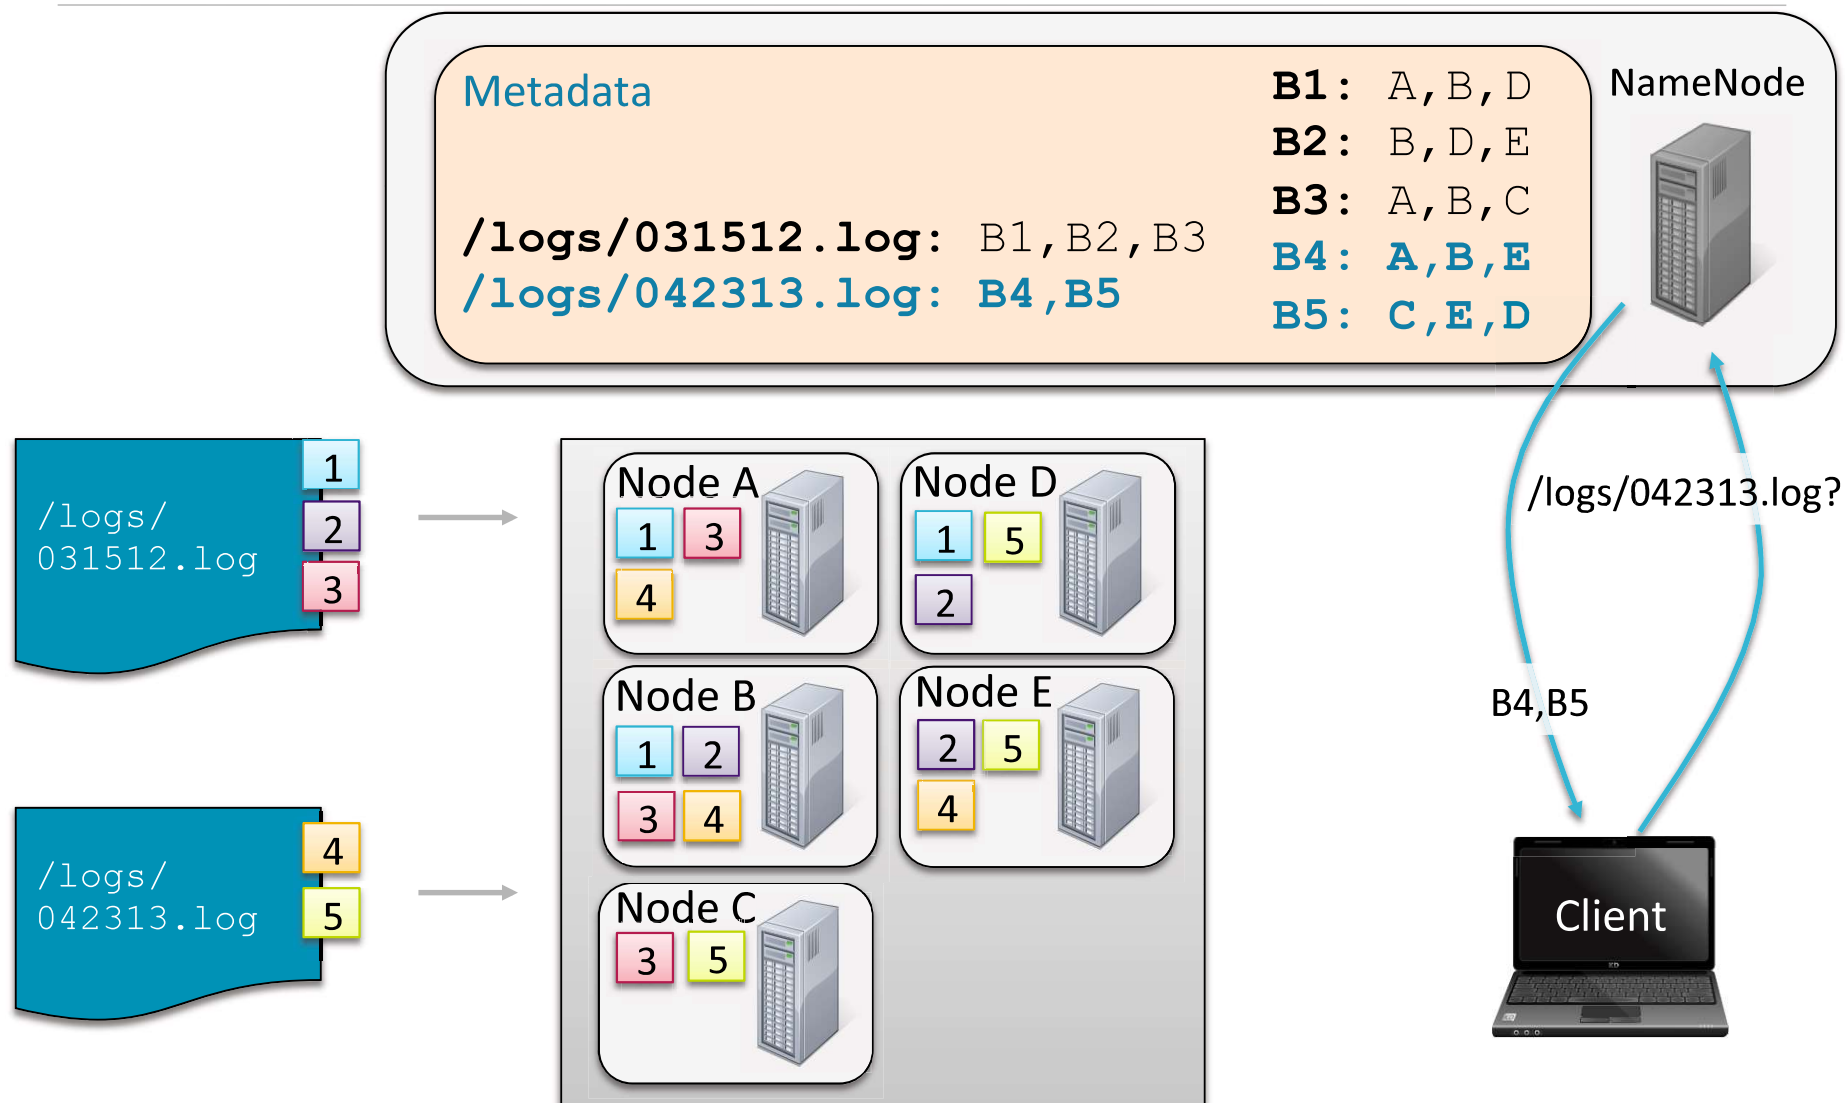
B4: A, B, E

B5: C, E, D

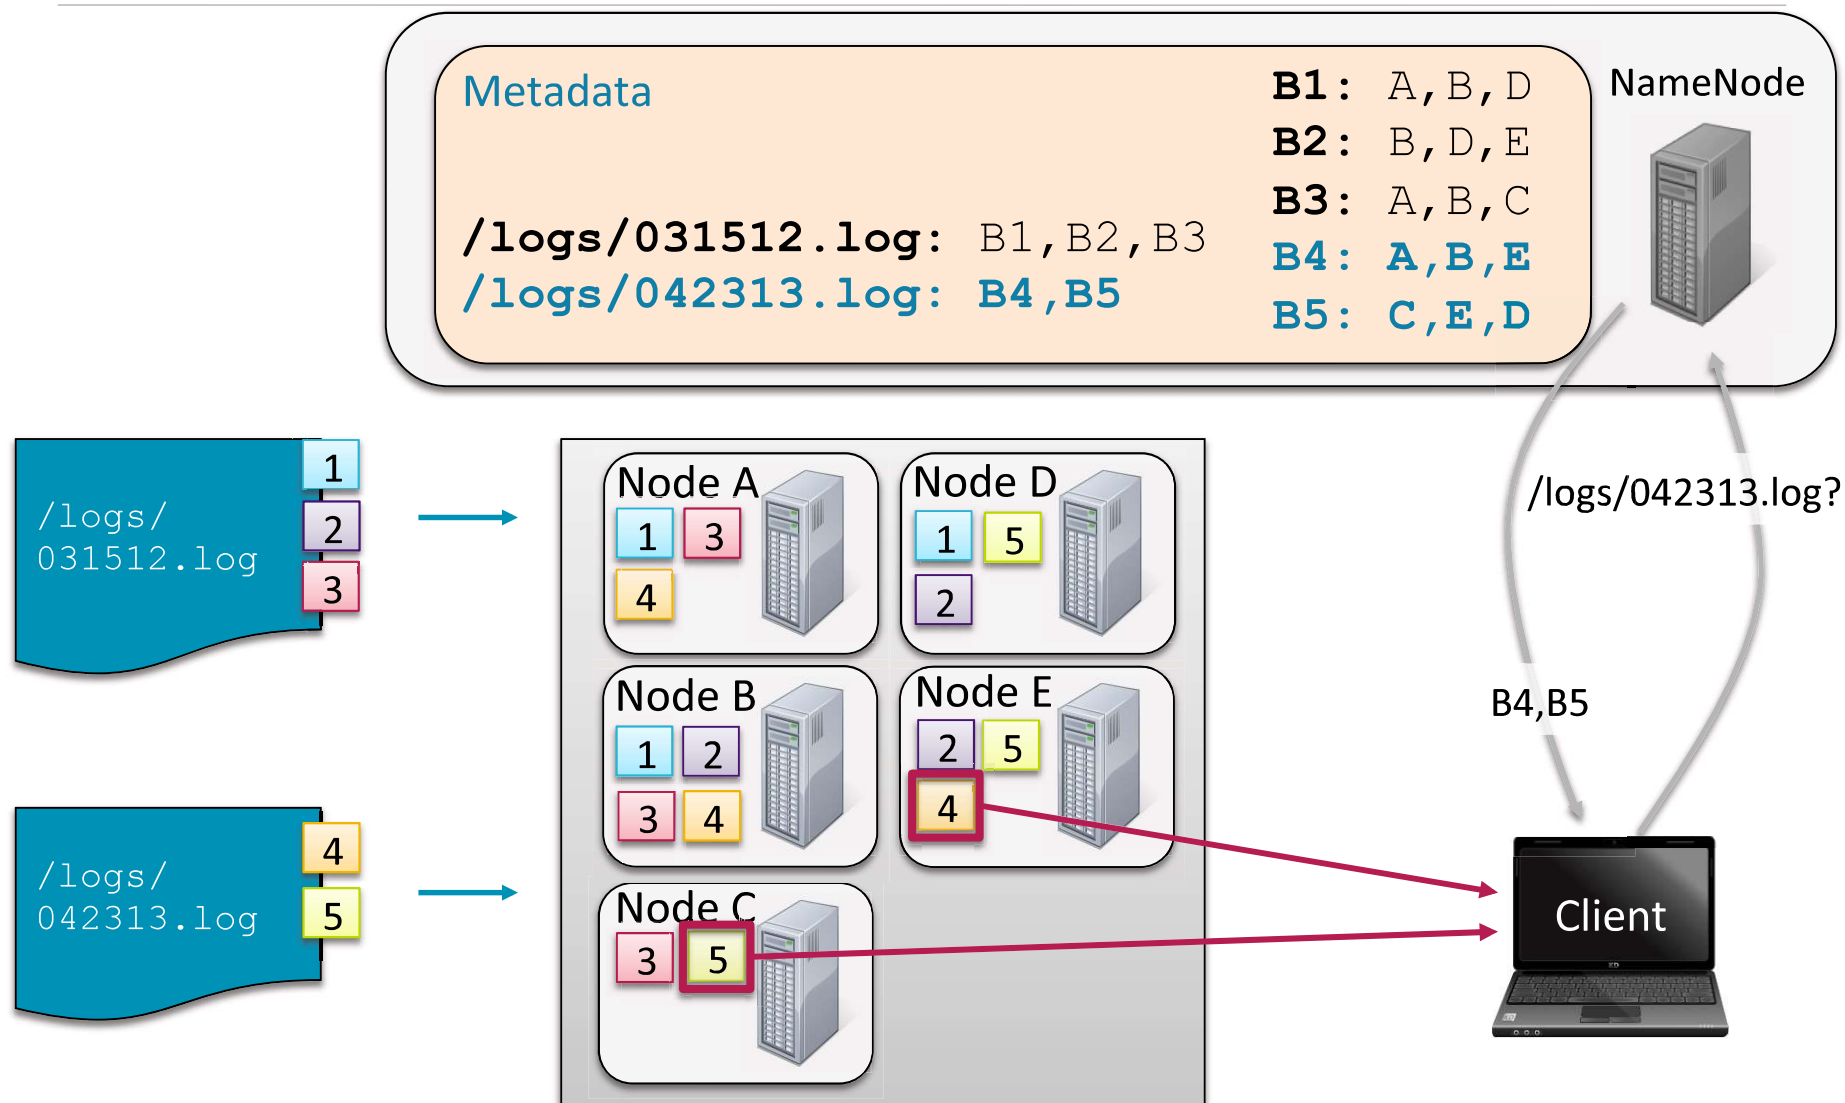
NameNode



Example: Storing and Retrieving Files (3)



Example: Storing and Retrieving Files (4)





HDFS NameNode Availability

- **The NameNode daemon must be running at all times**
 - If the NameNode stops, the cluster becomes inaccessible

- **HDFS is typically set up for High Availability**



- Two NameNodes: Active and Standby



High availability

- **Small clusters may use ‘Classic mode’**

- One NameNode
 - One “helper” node called the Secondary NameNode
 - Bookkeeping, not backup



Classical mode: Not highly available; namenode is

Options for Accessing HDFS

■ From the command line

– FsShell:

```
$ hdfs dfs
```

■ In Spark

– By URI, e.g.

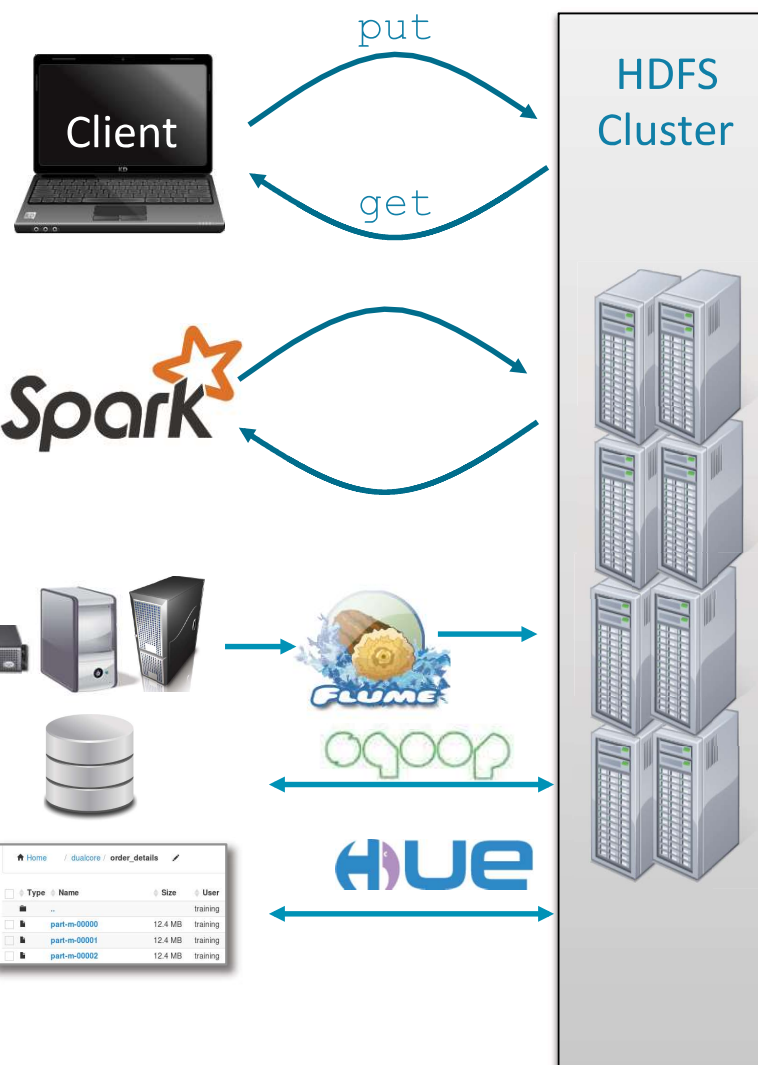
```
hdfs://nnhost:port/file...
```

■ Other programs

– Java API

– Used by Hadoop MapReduce, Impala, Hue, Sqoop, Flume, etc.

– RESTful interface



terminal com

Hadoop API

HDFS Command Line Examples (1)

- Copy file `foo.txt` from local disk to the user's directory in HDFS

```
$ hdfs dfs -put foo.txt foo.txt
```

– This will copy the file to `/user/username/foo.txt`

- Get a directory listing of the user's home directory in HDFS

```
$ hdfs dfs -ls
```

- Get a directory listing of the HDFS root directory

```
$ hdfs dfs -ls /
```

HDFS Command Line Examples (2)

- Display the contents of the HDFS file `/user/fred/bar.txt`

```
$ hdfs dfs -cat /user/fred/bar.txt
```

- Copy that file to the local disk, named as `baz.txt`

```
$ hdfs dfs -get /user/fred/bar.txt baz.txt
```

- Create a directory called `input` under the user's home directory

```
$ hdfs dfs -mkdir input
```

Note: `copyFromLocal` is a synonym for `put`; `copyToLocal` is a synonym for `get`

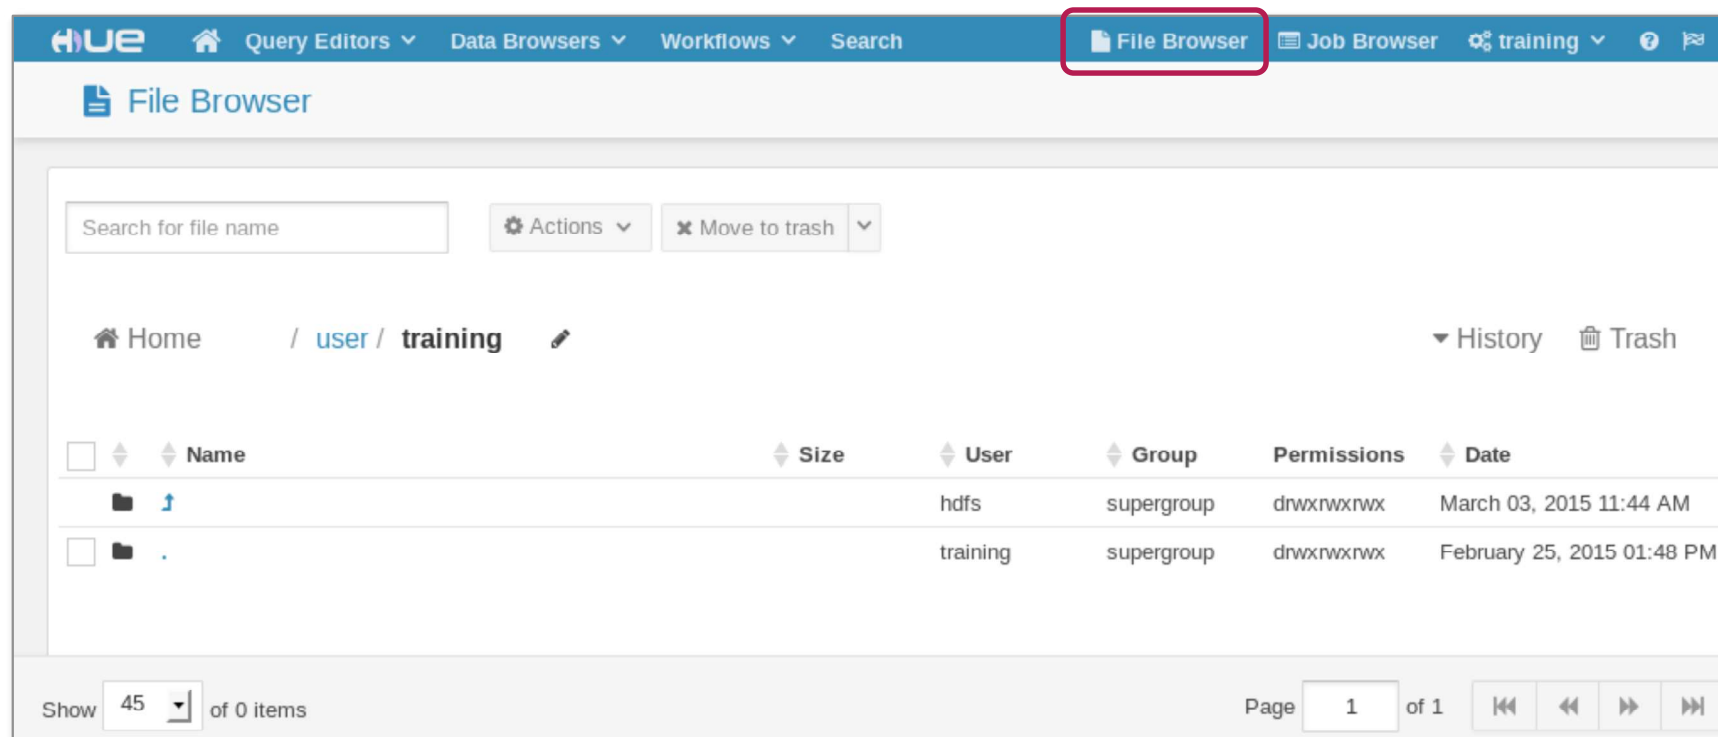
HDFS Command Line Examples (3)

- Delete the directory `input_old` and all its contents

```
$ hdfs dfs -rm -r input_old
```

The Hue HDFS File Browser

- **The File Browser in Hue lets you view and manage your HDFS directories and files**
 - Create, move, rename, modify, upload, download and delete directories and files
 - View file contents



HDFS Recommendations

- **HDFS is a repository for all your data**
 - Structure and organize carefully!
- **Best practices include**
 - Define a standard directory structure
 - Include separate locations for staging data
- **Example organization**
 - **/user/...** – data and configuration belonging only to a single user
 - **/etl** – Work in progress in Extract/Transform/Load stage
 - **/tmp** – Temporary generated data shared between users
 - **/data** – Data sets that are processed and available across the organization for analysis
 - **/app** – Non-data files such as configuration, JAR files, SQL files, etc.